

train_roberta_heavy

April 13, 2026

1 AGL — RoBERTa Fine-Tuning (Heavy Workload)

Adaptive Guardrail Layer (AGL) — Binary Prompt Safety Classifier

Fine-tunes `roberta-base` on the AGL binary dataset with **increased workload** for more training data. This notebook is designed for **local execution** (NVIDIA GPU with CUDA).

Data source: `data/processed/dataset_feature_engineered.csv`

Differences from baseline (train_roberta_local.ipynb): - Max samples per class: 25,000 (vs 15,000 baseline) - Epochs: 5 (vs 3 baseline) - Batch size: 48 (vs 32 baseline) - Early stopping patience: 3 (vs 2 baseline)

```
[20]: # 1. Setup
import os
import sys

# Set working directory to project root (parent of notebooks/)
PROJECT_ROOT = os.path.abspath(os.path.join(os.getcwd(), '..'))
if os.path.basename(os.getcwd()) == 'notebooks':
    os.chdir(PROJECT_ROOT)
elif not os.path.exists('src/config.py'):
    if os.path.exists('../src/config.py'):
        os.chdir('..')

if PROJECT_ROOT not in sys.path:
    sys.path.insert(0, os.getcwd())

print(f"Working directory: {os.getcwd()}")
assert os.path.exists('src/config.py'), "Not in project root - run from the_
↳notebooks/ or project root directory"
```

Working directory: `/Users/apadin/Desktop/Capstone/aai-590-group-8-capstone`

```
[21]: # 2. Check environment
import torch

print(f"Python: {sys.version}")
print(f"PyTorch: {torch.__version__}")
```

```

if hasattr(torch.backends, 'mps') and torch.backends.mps.is_available():
    DEVICE = "mps"
    print(f"Device: MPS (Apple Silicon)")
elif torch.cuda.is_available():
    DEVICE = "cuda"
    print(f"Device: CUDA - {torch.cuda.get_device_name(0)}")
else:
    DEVICE = "cpu"
    print(f"Device: CPU (training will be slow but it works)")

# Verify data exists
data_csv = 'data/processed/dataset_feature_engineered.csv'
assert os.path.exists(data_csv), f"Data not found at {data_csv}. Run the data_
pipeline notebooks first."
print(f"\nData: {data_csv} ")

```

Python: 3.12.4 | packaged by Anaconda, Inc. | (main, Jun 18 2024, 10:07:17)
[Clang 14.0.6]
PyTorch: 2.5.1
Device: MPS (Apple Silicon)

Data: data/processed/dataset_feature_engineered.csv

1.1 Step 3: Build Dataset (Heavy)

Load CSV, deduplicate, balance (25k per class), and split into train/val/test parquets.

```

[23]: # 3. Build dataset (HEAVY - 25k per class)
# Override the max samples per class for heavier training
import pandas as pd
from src.config import MAX_SAMPLES_PER_CLASS

# Temporarily override for heavier dataset
original_max = MAX_SAMPLES_PER_CLASS
import src.config as cfg
cfg.MAX_SAMPLES_PER_CLASS = 25000 # Heavy: 25k per class (50k total)

from src.data.build_dataset import build_dataset

splits = build_dataset(csv_path='data/processed/dataset_feature_engineered.csv')

for name, df in splits.items():
    print(f" {name}: {len(df)} samples")

# Restore original
cfg.MAX_SAMPLES_PER_CLASS = original_max

```

```
=====
Building AGL dataset (binary classification)
=====
```

```
Loaded 104755 samples from data/processed/dataset_feature_engineered.csv
Label distribution: {0: 55108, 1: 49647}
Dedup removed 5665 duplicates
```

```
After dedup: 99090 samples
```

```
After balancing: 50000 samples
```

```
Label distribution:
```

```
label_name
```

```
Benign      25000
```

```
Malicious   25000
```

```
Saved train: 34999 samples -> /Users/apadin/Desktop/Capstone/aai-590-group-8-
capstone/data/processed/train.parquet
```

```
Saved val: 7500 samples -> /Users/apadin/Desktop/Capstone/aai-590-group-8-
capstone/data/processed/val.parquet
```

```
Saved test: 7501 samples -> /Users/apadin/Desktop/Capstone/aai-590-group-8-
capstone/data/processed/test.parquet
```

```
Saved metadata -> /Users/apadin/Desktop/Capstone/aai-590-group-8-
capstone/data/processed/dataset_metadata.json
```

```
train: 34999 samples
```

```
val: 7500 samples
```

```
test: 7501 samples
```

```
[24]: # 3b. Verify dataset
import pandas as pd
from src.config import ID2LABEL

for split in ['train', 'val', 'test']:
    df = pd.read_parquet(f'data/processed/{split}.parquet')
    print(f"\n{split}: {len(df)} samples")
    label_counts = df['label'].value_counts()
    for label_id, count in label_counts.items():
        print(f" {ID2LABEL[label_id]}: {count}")
```

```
train: 34999 samples
```

```
Malicious: 17500
```

```
Benign: 17499
```

```
val: 7500 samples
```

```
Malicious: 3750
```

```
Benign: 3750
```

```
test: 7501 samples
  Benign: 3751
  Malicious: 3750
```

1.2 Step 4: Fine-Tune RoBERTa (Heavy)

Train the binary classifier with **heavier workload**: - LR: 2e-5, warmup: 10%, weight decay: 0.01
- **Batch: 48** (vs 32 baseline) - **Epochs: 5** (vs 3 baseline) - **Early stopping patience: 3** (vs 2 baseline) - Max sequence length: 128 - 2 classes: Benign (0), Malicious (1)

```
[ ]: # 4. Train RoBERTa classifier (HEAVY)
# Override training parameters for heavier workload
import src.config as cfg

# Save original values
orig_batch = cfg.TRAIN_BATCH_SIZE
orig_epochs = cfg.NUM_EPOCHS
orig_patience = cfg.EARLY_STOPPING_PATIENCE

# Apply heavy settings
cfg.TRAIN_BATCH_SIZE = 48      # Larger batch
cfg.NUM_EPOCHS = 5             # More epochs
cfg.EARLY_STOPPING_PATIENCE = 3 # More patience

from src.training.train import train_classifier

best_path = train_classifier()
print(f"\nBest model saved to: {best_path}")

# Restore original
cfg.TRAIN_BATCH_SIZE = orig_batch
cfg.NUM_EPOCHS = orig_epochs
cfg.EARLY_STOPPING_PATIENCE = orig_patience
```

```
tokenizer_config.json:  0%|          | 0.00/25.0 [00:00<?, ?B/s]
config.json:  0%|          | 0.00/481 [00:00<?, ?B/s]
vocab.json: 0.00B [00:00, ?B/s]
merges.txt: 0.00B [00:00, ?B/s]
tokenizer.json: 0.00B [00:00, ?B/s]
Map:  0%|          | 0/34999 [00:00<?, ? examples/s]
Tokenized train: 34999 samples
Map:  0%|          | 0/7500 [00:00<?, ? examples/s]
Tokenized val: 7500 samples
Map:  0%|          | 0/7501 [00:00<?, ? examples/s]
```

Tokenized test: 7501 samples

Class weights: [1.0000285730613179, 0.9999714285714286]

model.safetensors: 0%| | 0.00/499M [00:00<?, ?B/s]

Some weights of RobertaForSequenceClassification were not initialized from the model checkpoint at roberta-base and are newly initialized:

['classifier.dense.bias', 'classifier.dense.weight', 'classifier.out_proj.bias', 'classifier.out_proj.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Starting training: 5 epochs, lr=2e-05

<IPython.core.display.HTML object>

1.3 Step 5: Evaluate

Run the full evaluation suite: keyword baseline, TF-IDF/SVM baseline, RoBERTa classifier.

```
[ ]: # 6. Run evaluation
import numpy as np
import pandas as pd
from src.config import MODELS_DIR, PROCESSED_DIR, RESULTS_DIR
from src.evaluation.metrics import evaluate_predictions, save_results
from src.evaluation.baselines import keyword_blocklist_baseline, tfidf_svm_baseline

# Load test data
train_df = pd.read_parquet(PROCESSED_DIR / 'train.parquet')
test_df = pd.read_parquet(PROCESSED_DIR / 'test.parquet')
y_true = test_df['label'].values

print("=" * 60)
print("Running evaluation suite (HEAVY model)")
print("=" * 60)

all_results = {}

# Baseline 1: Keyword blocklist
print("\n[1/3] Keyword blocklist baseline...")
kw_preds = keyword_blocklist_baseline(test_df)
kw_results = evaluate_predictions(y_true, kw_preds)
all_results['keyword_blocklist'] = kw_results
print(f" Macro-F1: {kw_results['macro_f1']:.4f}")

# Baseline 2: TF-IDF + SVM
print("\n[2/3] TF-IDF + LinearSVM baseline...")
svm_preds, _ = tfidf_svm_baseline(train_df, test_df)
```

```

svm_results = evaluate_predictions(y_true, svm_preds)
all_results['tfidf_svm'] = svm_results
print(f"  Macro-F1: {svm_results['macro_f1']:.4f}")

# RoBERTa
checkpoint = MODELS_DIR / 'classifier' / 'best'
if checkpoint.exists():
    print("\n[3/3] RoBERTa classifier (HEAVY)...")
    from src.models.agl_pipeline import AGLPipeline
    pipeline = AGLPipeline.from_checkpoint(checkpoint)
    roberta_preds = np.array([pipeline.predict(t).label_id for t in
    ↪test_df['text']])
    roberta_results = evaluate_predictions(y_true, roberta_preds)
    all_results['roberta_heavy'] = roberta_results
    print(f"  Macro-F1: {roberta_results['macro_f1']:.4f}")
else:
    print(f"\n[3/3] Skipping RoBERTa - no checkpoint at {checkpoint}")

save_results(all_results, 'evaluation_results_heavy', RESULTS_DIR)

print("\n" + "=" * 60)
print("Summary - Macro-F1 Scores:")
print("=" * 60)
for method, res in all_results.items():
    print(f"  {method:25s} {res['macro_f1']:.4f}")

```

1.4 Step 6: Visualizations

Generate figures: confusion matrix, F1 comparison, ROC curve.

```

[ ]: # 7. Generate visualizations
import numpy as np
import pandas as pd
import torch
from src.config import RESULTS_DIR, FIGURES_DIR, PROCESSED_DIR, MODELS_DIR,
    ↪MAX_SEQ_LENGTH
from src.evaluation.visualizations import (
    plot_confusion_matrix, plot_f1_comparison, plot_roc_curves
)
from src.evaluation.metrics import evaluate_predictions, benchmark_latency,
    ↪save_results
from src.evaluation.baselines import keyword_blocklist_baseline,
    ↪tfidf_svm_baseline
from src.models.agl_pipeline import AGLPipeline
from transformers import AutoTokenizer

FIGURES_DIR.mkdir(parents=True, exist_ok=True)

```

```

# Load test data
test_df = pd.read_parquet(PROCESSED_DIR / 'test.parquet')
train_df = pd.read_parquet(PROCESSED_DIR / 'train.parquet')
y_true = test_df['label'].values

# Baselines
kw_preds = keyword_blocklist_baseline(test_df)
svm_preds, _ = tfidf_svm_baseline(train_df, test_df)

# RoBERTa predictions + probabilities (for ROC-AUC)
checkpoint = MODELS_DIR / 'classifier' / 'best'
pipeline = AGLPipeline.from_checkpoint(checkpoint)
tokenizer = AutoTokenizer.from_pretrained(str(checkpoint),
    ↪local_files_only=True)

roberta_preds = []
roberta_probs = []
for text in test_df['text']:
    pred = pipeline.predict(text)
    roberta_preds.append(pred.label_id)
    inputs = tokenizer(text, truncation=True, padding='max_length',
        max_length=MAX_SEQ_LENGTH, return_tensors='pt')
    inputs = {k: v.to(pipeline.device) for k, v in inputs.items()}
    with torch.no_grad():
        logits = pipeline.model(**inputs).logits
        probs = torch.softmax(logits, dim=-1).cpu().numpy()[0]
    roberta_probs.append(probs)

roberta_preds = np.array(roberta_preds)
roberta_probs = np.array(roberta_probs)

# Confusion matrix
plot_confusion_matrix(y_true, roberta_preds, title='RoBERTa (Heavy)',
    save_path=FIGURES_DIR / 'confusion_matrix_roberta_heavy.
    ↪png')

# F1 comparison
from sklearn.metrics import f1_score
f1_results = {
    'Keyword Blocklist': f1_score(y_true, kw_preds, average='macro',
    ↪zero_division=0),
    'TF-IDF + SVM': f1_score(y_true, svm_preds, average='macro',
    ↪zero_division=0),
    'RoBERTa (Heavy)': f1_score(y_true, roberta_preds, average='macro',
    ↪zero_division=0),
}

```

```

plot_f1_comparison(f1_results, save_path=FIGURES_DIR / 'f1_comparison_heavy.
    ↪png')

# ROC Curve
plot_roc_curves(y_true, roberta_probs, title='RoBERTa Heavy ROC Curve',
    save_path=FIGURES_DIR / 'roc_curve_roberta_heavy.png')

# Full evaluation with AUC
results = evaluate_predictions(y_true, roberta_preds, y_prob=roberta_probs)
print(f"\nRoBERTa (Heavy) Results:")
print(f"  Macro-F1:  {results['macro_f1']:.4f}")
print(f"  ROC-AUC:    {results.get('roc_auc', 'N/A')}")
print(f"  PR-AUC:     {results.get('pr_auc', 'N/A')}")

# Latency
latency = benchmark_latency(pipeline, test_df['text'].head(50).tolist(),
    ↪n_runs=2)
print(f"\nLatency: mean={latency['mean_ms']:.1f}ms, p99={latency['p99_ms']:.
    ↪1f}ms")

# Save results JSON
all_results = {
    'roberta_heavy': {**results, 'latency': latency}
}
save_results(all_results, 'roberta_heavy_evaluation', RESULTS_DIR)
print('\nAll figures saved to:', FIGURES_DIR)

```

1.5 Step 7: Done

Model checkpoint saved to `models/classifier/best/`, results to `results/`.

```

[ ]: # 7. Summary
from pathlib import Path

print("Heavy Training Complete!")
print("Artifacts saved locally:")
for p in ['models/classifier/best', 'results']:
    exists = " " if Path(p).exists() else " "
    print(f"  {exists} {p}")

print("\nComparison with baseline:")
print("  Baseline (train_roberta_local.ipynb): 15k samples/class, 3 epochs,
    ↪batch 32")
print("  Heavy (this notebook): 25k samples/class, 5 epochs, batch 48")

```